

CECS 274: Data Structures

Hash Tables

Oscar Morales Ponce

California State University Long Beach

Example


- ▶ We can check if the following expression is valid using a Stack (Array-based, Linked-List based)

$$((a + b) * (c * d) + (a + (b * (c/d))))$$

- ▶ How can we replace the terms with the actual values efficiently:
 $a = 5413.13, b = 243.12, c = 4212.12, d = 421.33$
- ▶ Suppose you have a list of several millions of unique items. How to search efficiently when the name of the item is known.

Hash Table

- ▶ Hash tables are an efficient method of storing a small number, n , of integers from a large range $U = \{0, \dots, 2^w - 1\}$.
- ▶ The term *hash table* includes a broad range of data structures.
- ▶ The *hash value* of a data item x , denoted $\text{hash}(x)$ is a value in the range $\{0, \dots, 2^d - 1\}$ for some $d \geq 0$



Input space

Hash space



ChainedHashTable: The Basics


- ▶ ChainedHashTable

- ▶ State variables:

- ▶ n : Number of elements in a .
 - ▶ d : Determines the size of the array
 - ▶ t : Array of lists of length 2^d
 - ▶ z : A random odd integer

- ▶ Invariant: $n \leq \text{length}(t) = 2^d < 3n$ and

For a pair $(key, value)$, $value$ is uniquely stored in the list at $t[\text{hash}(key)]$



ChainedHashTable: *initialize()*

- *initialize()*: Initialize the state variables

```
initialize()  
   $d = 1$   
   $t \leftarrow \text{alloc\_table}(2^d)$   
   $z \leftarrow \text{random\_odd\_int}()$   
   $n \leftarrow 0$ 
```

ChainedHashTable: $add(key, value)$

- ▶ $add(x, v)$: Add the tuple (x, v) in the table if it does not exist

$add(key, value)$

Check Preconditions

Check that the invariant holds

Insert $Node(key, value)$ in the list $t[hash(key)]$

Increment n by one

ChainedHashTable: $add(key, value)$

```
add(key, value)  
  if find(key)  $\neq$  nil then return false  
  if  $n = \text{length}(t)$  then resize()  
   $t[\text{hash}(key)].\text{add}(0, \text{Node}(key, value))$   
   $n \leftarrow n + 1$   
  return true
```

- ▶ The cost of growing is only constant when amortized over a sequence of insertions
- ▶ Appending $key, value$ to the list $t[\text{hash}(key)]$ takes only constant time.

ChainedHashTable: $remove(key)$

- ▶ $remove(key)$: Remove the element (key, \cdot) from the table if it exists

$remove(key)$

Check Preconditions

Check if the value key exists in $t[hash(key)]$

Check that the Invariant holds

ChainedHashTable: *remove(key)*

```
remove(key)  
  for i in 0, 1, ..., t[hash(key)].size() - 1  
    if t[hash(key)].get(i).key = key then  
      t[hash(key)].remove(i)  
      n ← n - 1  
      if length(a) ≥ 3 · n then resize()  
      return true  
return false
```

- ▶ This takes $O(n_{\text{hash}(x)})$ time, where n_i denotes the length of the list stored at $t[\text{hash}(key)]$.

ChainedHashTable: *resize()* (Discussion Activity Write the Steps)

- *resize()*: Resize the table t such that $n \leq \text{length}(t) < 3n$

resize()

ChainedHashTable: *resize()*

```
resize()  
  if  $n = \text{length}(t)$  then  $d + 1$   
  else  $d - 1$   
   $a \leftarrow \text{alloc\_table}(2^d)$   
  for  $j$  in  $0, 1, \dots, \text{length}(t) - 1$   
    for  $i$  in  $0, 1, \dots, t[j].\text{size}() - 1$   
       $a[\text{hash}(t[j].\text{get}(i).\text{key})].\text{add}(t[j].\text{get}(i))$   
   $t = a$ 
```

ChainedHashTable: $find(key)$

- ▶ $find(key)$: Returns the node if key exists and None otherwise

$find(key)$

Check Preconditions

Check if key exist in $t[hash(key)]$ and return the node.
return nil

ChainedHashTable: *find(key)*

```
find(key)  
  for i in 0, 1, .., t[hash(key)].size() - 1  
    if t[hash(key)].get(i).key = key then  
      return t[hash(key)].get(i)  
  return nil
```

- ▶ This takes time proportional to the length of the list $t[\text{hash}(x)]$.

ChainedHashTable: Performance

- ▶ The performance of a hash table depends critically on the choice of the hash function.
- ▶ A good hash function will spread the elements evenly among the $\text{length}(t)$ lists, so that the expected size of the list $t[\text{hash}(x)]$ is $O(n/\text{length}(t)) = O(1)$.
- ▶ On the other hand, a bad hash function will hash all values to the same table location, in which case the size of the list $t[\text{hash}(key)]$ will be n .

Multiplicative Hashing

- ▶ It uses the *div* operator, which calculates the integral part of a quotient, while discarding the remainder.
- ▶ Formally, for any integers $a \geq 0$ and $b \geq 1$, $a \text{ div } b = \lfloor a/b \rfloor$.
- ▶ We use a hash table of size 2^d for some integer d (called the *dimension*).
- ▶ The formula for hashing an integer $key \in \{0, \dots, 2^w - 1\}$ is

$$\text{hash}(key) = ((z \cdot \text{hashCode}(key) \bmod 2^w) \text{ div } 2^{w-d})$$

- ▶ Here, z is a randomly chosen *odd* integer in $\{1, \dots, 2^w - 1\}$.
- ▶ By default operations on integers are already done modulo 2^w where w is the number of bits in an integer

Multiplicative Hashing

- ▶ Integer division by 2^{w-d} is equivalent to dropping the rightmost $w - d$ bits in a binary representation
- ▶ It is implemented by shifting the bits right by $w - d$ using the $>>$ operator

```
hash(key)  
    return ((z · hashCode(key)) mod  $2^w$ )  $>>$  ( $w - d$ )
```


Multiplicative Hashing

► Python

```
def _hash(self, x):  
    return ((self.z * hash(x)) % (1<<w)) >> (w-self.d)
```

► Java

```
int hash(Object x) {  
    return (z * x.hashCode()) >>> (w-d);  
}
```

► C++

```
int hash(T x) {  
    return ((unsigned)(z * std::hash<std::string>{}(x)))  
           >> (w-d);  
}
```

Universal Hashing

Lemma

Let x and y be any two values in $\{0, \dots, 2^w - 1\}$ with $x \neq y$.
Then $\Pr\{\text{hash}(x) = \text{hash}(y)\} \leq 2/2^d$.

Proof

Observe that the highest-order d bits of $zx \bmod 2^w$ and dz bits of z are the same. Therefore, the highest-order d bits in the binary representation of $z(x - y) \bmod 2^w$ are either all 0's or all 1's. That is,

$$z(x - y) \bmod 2^w = (\underbrace{0, \dots, 0}_d, \underbrace{\star, \dots, \star}_{w-d})_2 \quad (1)$$

when $zx \bmod 2^w > zy \bmod 2^w$ or

$$z(x - y) \bmod 2^w = (\underbrace{1, \dots, 1}_d, \underbrace{\star, \dots, \star}_{w-d})_2 . \quad (2)$$

when $zx \bmod 2^w < zy \bmod 2^w$.

Universal Hashing (Proof 2/3)

Therefore, we only have to bound the probability that

$z(x - y) \bmod 2^w$ looks like (2) or (3).

Let q be the unique odd integer such that $(x - y) \bmod 2^w = q2^r$ for some integer $r \geq 0$. The binary representation of $zq \bmod 2^w$ has $w - 1$ random bits, followed by a 1:

$$zq \bmod 2^w = (\underbrace{b_{w-1}, \dots, b_1}_{w-1}, 1)_2$$

Therefore, the binary representation of

$z(x - y) \bmod 2^w = zq2^r \bmod 2^w$ has $w - r - 1$ random bits, followed by a 1, followed by r 0's:

$$z(x - y) \bmod 2^w = zq2^r \bmod 2^w = (\underbrace{b_{w-r-1}, \dots, b_1}_{w-r-1}, 1, \underbrace{0, 0, \dots, 0}_r)_2$$

Universal Hashing (Proof 3/3)

- If $r > w - d$, then the d higher order bits of $z(x - y) \bmod 2^w$ contain both 0's and 1's, so the probability that $z(x - y) \bmod 2^w$ looks like (2) or (3) is 0.
- If $r = w - d$, then the probability of looking like (0) is 0, but the probability of looking like (1) is $1/2^{d-1} = 2/2^d$ (since we must have $b_1, \dots, b_{d-1} = 1, \dots, 1$).
- If $r < w - d$, then we must have $b_{w-r-1}, \dots, b_{w-r-d} = 0, \dots, 0$ or $b_{w-r-1}, \dots, b_{w-r-d} = 1, \dots, 1$. The probability of each of these cases is $1/2^d$ and they are mutually exclusive, so the probability of either of these cases is $2/2^d$.

Expected Value and Probability

- ▶ For a discrete random variable X taking on values in some countable universe U , the expected value of X , denoted by $E[X]$, is given by

$$E[X] = \sum_{x \in U} x \cdot \Pr\{X = x\}$$

Where $\Pr\{\mathcal{E}\}$ denotes the probability that the event \mathcal{E} occurs.

- ▶ Linearity of expectation. For any two random variables X and Y ,

$$E[X + Y] = E[X] + E[Y]$$

More generally, for any random variables X_1, \dots, X_k ,

$$E\left[\sum_{i=1}^k X_i\right] = \sum_{i=1}^k E[X_i]$$

Hash Analysis

Lemma

For any unique key x , the expected length of the list $t[\text{hash}(x)]$ is at most 2.

Proof Let S be the set of elements stored in the hash table that are not equal to x . For an element $y \in S$, define the indicator variable

$$I_y = \begin{cases} 1 & \text{if } \text{hash}(x) = \text{hash}(y) \\ 0 & \text{otherwise} \end{cases}$$

Proof

By the previous lemma, $E[I_y] \leq 2/2^d = 2/\text{length}(t)$.
The expected length of the list $t[\text{hash}(x)]$ is given by

$$\begin{aligned} E[t[\text{hash}(x)].\text{size}()] &= E\left[\sum_{y \in S} I_y\right] \\ &\leq \sum_{y \in S} 2/\text{length}(t) \\ &\leq \sum_{y \in S} 2/n \\ &= 2 \end{aligned}$$

Summary of ChainedHashTable

Theorem

The expected running time of $\text{add}(x)$, $\text{remove}(x)$ and $\text{find}(x)$ of ChainHashTable is $O(1)$.